PAVEL CHERKASHIN

TASTE LABS — AI ENGINEER, RL & EVALS (TECHNICAL STAFF)

San Francisco, CA • In-Person SF
pavel@mindrock.capital • linkedin.com/in/cherkashin
github.com/pashamango • @pashamango

ROLE SEAT & POSTURE: Evals-focused systems architect applying rigorous empirical evaluation harnesses, decomposition grading, and multi-model taste rubrics to end "AI slop" in subjective domains. Targeting **AI Engineer, RL & Evals** at Taste Labs to take end-to-end ownership of **grading methods & rubrics, task design for taste, agent harnesses & context layers, and deterministic evaluation gates**. *Transparent alignment:* deep strength in evals harnesses, rubrics decomposition, agent context layers, backend infrastructure, and data indexing; transparently recognizing scalable RL cluster infra and distributed training pipelines as an active growth edge rather than core depth.

CORE COMPETENCIES & TASTE EVALS ARCHITECTURE

Subjective Domain Evals & Taste Rubrics
- Decomposing design, tone & aesthetics into granular, testable rubrics
- Multi-model LLM-as-judge scoring & pairwise Bradley-Terry Elo ranking
- Held-out private verifiers to prevent reward hacking & benchmark memorization
- Empirical failure-mode taxonomies & programmatic anti-slop boundary assertions

Backend Infrastructure & Data Pipelines
- High-throughput distributed backends & event-driven architecture
- Deterministic test harnesses, isolated worktrees & CI/CD verification gates
- Tamper-evident cryptographic manifests (SHA-256) & zero-trust verification
- Large-scale dataset ingestion, feature extraction & telemetry streaming

Agent Harnesses & Context Layers
- Autonomous agent runtime orchestration & daemon architecture (*fleetd*)
- Dynamic context layering, memory injection & prompt state management
- Model Context Protocol (MCP) tool integration & sandbox isolation
- Strict concurrency gating (WIP=1), state auditability & crash recovery

Technical Pedagogy & Creative Systems
- Published textbook author: *CRM Strategy* (INTUIT/BINOM, Moscow, 384 pp.)
- Turning fuzzy human taste and design guidelines into structured engineering specs
- Clear public benchmark documentation, research notes & anti-gaming disclosures
- Enterprise solution architecture, executive communication & partner alignment

SYSTEMS & LEADERSHIP EXPERIENCE

Mindrock Autonomous Agent Fleet — *Principal Systems Architect & Harness Builder* 2024 – Present | San Francisco, CA
- **Agent Runtime Architecture:** Architect of *fleetd*, a monolithic daemon in Rust managing multi-agent concurrency, slot leasing (WIP=1), deterministic task dispatch, and autonomous backlog vacuum loops across distributed macOS nodes.
- **Deterministic Evals & QA Gates:** Built Protocol Phi (Φ-0.1), establishing structured, tamper-evident handoffs and cross-agent evaluation loops between builder and independent QA agents, enforcing programmatic assertions before human gate review.
- **Multi-Model Verification & Anti-Slop:** Engineered automated evaluation harnesses combining geometric layout assertions, multi-model consensus, and LLM-as-judge scoring, eliminating subjective drift, visual collisions, and synthetic hallucination.
- **Model Context Protocol (MCP):** Integrated native MCP servers with lazy schema resolution, sandboxed execution, and token-efficient tool streaming.

Vestor.in & Technology Advisory — *Founding Partner & Systems Architect* 2018 – 2024 | San Francisco, CA
- **Platform Architecture:** Designed high-concurrency event routing and data platforms for technology investments; evaluated early-stage AI architectures and product-market fit across generative AI, autonomous agents, and enterprise automation.
- **Startup Coaching & Systems Strategy:** Advised founders on technical roadmap design, automated verification pipelines, and scalable backend infrastructure.

Sputnik Labs — *Founder & Chief Executive Officer* (Acquired by Technoserv, 2008) 2001 – 2008 | Moscow & European Hubs
- **Founding & Enterprise Scale:** Founded and scaled the pioneering enterprise CRM and workflow systems integration firm in Eastern Europe; built engineering organization to 500+ employees and led successful acquisition by Technoserv.
- **Customer Rubrics & Predictive Scoring:** Architected large-scale customer interaction rubrics, automated data indexing pipelines, and real-time behavioral scoring systems for tier-1 telecom and banking enterprises in strategic partnership with Siebel, Oracle, and Microsoft.

Siebel Systems & Microsoft — *Country Manager, Russia & CIS / Platform Strategy Director* 1999 – 2006 | Regional Executive Roles
- **Platform Leadership:** Directed enterprise software platforms, developer evangelism, and integration architectures for mission-critical enterprise systems, establishing regional technology leadership and partner integration frameworks.

EDUCATION & PUBLISHED AUTHOR

Lomonosov Moscow State University (MSU) — *B.S. in Computer Science in Global Ecology* 1995 | Moscow, RU
- Mathematical modeling, computational data structures, satellite telemetry analysis. Research fellowship at the International Institute for Applied Systems Analysis (IIASA, Laxenburg, Austria, 1994–1995).

Published University Textbook: *Customer Relationship Management Strategy* (*[redacted] (CRM)*)
- INTUIT / BINOM, Moscow, 2007, 384 pp., **ISBN 978-5-94774-643-3**. Cataloged in the Russian State Library (RSL Official Record: 01003185368). Definitive academic text on automated customer data modeling, enterprise rubrics, and systematic interaction design.